

6.046 Problem 2-2Collaborators: *none*

(a) The input of this problem are a key k and two B-trees T_1 and T_2 with same minimum degree parameter t and same height $h_1 = h_2$. All keys in T_1 are strictly smaller than k while all keys in T_2 are strictly larger than k . The out put should be another B-tree with minimum degree parameter t and contains exactly all keys in T_1 and T_2 plus k

Let the root node of T_1 and T_2 be r_1 and r_2 , respectively. Then we create a node r , whose keys are keys in r_1 , k , and keys in r_2 , in that order. The children of r are children of r_1 and children of r_2 , in their previous order, and the children of r_2 come after the children of r_1 . Now if $r.n \leq 2t - 1$, return r ; otherwise, create nodes r'_1 , r'_2 and r' , where the keys in r'_1 are the first half in r :

$$\{r.key_1, \dots, r.key_{\lfloor \frac{r.n-1}{2} \rfloor}\}$$

and the keys in r'_2 are

$$\{r.key_{\lfloor \frac{r.n+3}{2} \rfloor}, \dots, r.key_{[r.n]}\}$$

The children of r'_1 are $r.c_1, \dots, r'_1$ is $r.c_{\lfloor \frac{r.n+1}{2} \rfloor}$; the children of r'_2 are $r.c_{\lfloor \frac{r.n+3}{2} \rfloor}, \dots, r.c_{r.n}$. Let r' has only one key $r.key_{\lfloor \frac{r.n+1}{2} \rfloor}$, and $r'.c_1 = r'_1$ and $r'.c_2 = r'_2$. Then return r' .

The tree starting from r we create in this algorithm contains exactly all keys in T_1 , T_2 , and k , and as all its nodes except for the root node are from either T_1 or T_2 , the number of keys they contain is between $t - 1$ and $2t - 1$. As T_1 and T_2 has same height, r leads a B-tree except for that the $r.n$ might exceed $2t - 1$. If $r.n \leq 2t - 1$, r itself is qualified. However, if $r.n > 2t - 1$, as both $r_1.n \leq 2t - 1$ and $r_2.n \leq 2t - 1$, we have that $r.n \leq 4t - 1$. Thus both $r'_1.n$ and $r'_2.n$ are $\leq 2t - 1$ and $> t - 1$. As r' contains exactly all keys in r , and all other nodes in r' are from T_1 or T_2 , we might say that r' is a qualified tree.

Making r , r'_1 , r'_2 and r' all takes $O(t)$ time, and this time is considered constant.

(b) The input of this problem are a key k and two B-trees T_1 and T_2 with same minimum degree parameter t and have heights h_1 and h_2 , where $h_1 = h_2 + 1$. All keys in T_1 are strictly smaller than k while all keys in T_2 are strictly larger than k . The out put should be another B-tree with minimum degree parameter t and contains exactly all keys in T_1 and T_2 plus k

Let the roots of T_1 and T_2 are r_1 and r_2 respectively. Let the last child of r_1 be $r_1.lc$. Then $r_1.lc$ has the same height as r_2 . We do the same algorithm in (a) on $r_1.lc$ and r_2 , and assume the result is r_g . If the tree leading by r_g has height $h_1 - 1$, this indicates $r_g.n = r_1.lc + r_2 > t - 1$, and $r_g.n < 2t - 1$. We can replace $r_1.lc$ with r_g , and return r_1 .

On the other hand, if the tree leading by r_g has height h_1 , this indicates that $r_g.n = 1$. Add a key, which equals $r_g.key_1$, to r , and that key should be the largest key of r . Let the two nodes before and after that key be the two child of r_g (thus $r_1.lc$ is overwritten). Now return r_1 if $r_1.n \leq 2t - 1$; otherwise, do the same split that we did on r in (a), and return the tree we get.

According to (a), the r_g we get here is a qualified B-tree with all keys from the tree under $r_1.lc$ and T_2 . If it has height $h_1 - 1$, or h_2 , replacing $r_1.lc$ with r_g simply adds all keys in T_2 . As $r_g.n$ is within the range, r_1 is what we need. If r_g has height h_1 , the process we take also deletes $r_1.lc$ from the tree leading by r_1 and adds r_g , and the number of keys in all nodes are between $t - 1$ and $2t - 1$, except for r_1 , which might have $2t$ keys. This can be handled using the same approach we dealt with r in (a).

Getting r_g requires constant time, according to (a). Merging r_g with r_1 under both conditions also takes constant time, as we need at most add one key to r_1 and redirect two children. The possible splitting of r_1 still takes constant time, according to (a). Thus the total time is constant.

(c) The insertion takes an augmented B-tree and a key k not in the tree as input, returns another augmented B-tree that contains all keys in the input B-tree and k ; the deletion takes an augmented B-tree and a key k that is in the tree as input, returns a B-tree that contains all keys in the input tree except for k . All B-trees here have the same minimum degree t .

Both algorithms are similar to the algorithms for non augmented B-tree. We let $x.h$ be the height of the subtree below x . For the insertion, the "B-TREE-SPLIT-CHILD(x, i)" in CLRS 18.2 should be modified. We need to let $z.h = y.h$ after creating z , as the subtree below y will be split to two below y and z , respectively. Besides, for "B-TREE-INSERT(T, i)", we should let $s.h = r.h + 1$ when creating s .

As at most one s will be created during one insertion, and split at most h times, the time added is $O(h)$, and the total time is $O(h)$.

For deletion, like in CLRS 18.3, we first check whether the root node is empty; if empty, we

let its only child be the root instead. Then call a function "DELETE-NON-MINIMUM(x, k)", which tries to delete k from the subtree below x , on condition that k is in the subtree, and $x.n$ is not minimum (for non root is $t - 1$, root is 0). In "DELETE-NON-MINIMUM", we first check if k is a key in x : if true, check the child before and after k : if the number of keys in either doesn't reach the minimum, find the predecessor or successor of k , and run "DELETE-NON-MINIMUM(*child before* $k, k.predecessor$)" or "DELETE-NON-MINIMUM(*child after* $k, k.successor$)". Then replace k with $k.predecessor$ or $k.successor$. If both reaches the minimum ($t - 1$), merge the two children to one node, which has the same height as the two children, and keys from both and k . Then delete k from the new node.

If k is not a key in x , find the child $x.c_j$ of x that contains k . If $x.c_j.n > t - 1$, just call "DELETE-NON-MINIMUM($x.c_j, k$)"; otherwise, check $x.c_{j-1}.n$ and $x.c_{j+1}$. If either is $> t - 1$, add $x.key_j/x.key_{j+1}$ to $x.c_j$, move the last/first key of $x.c_{j-1}/x.c_{j+1}$ to $x.key_j/x.key_{j+1}$, move the last/first subtree of $x.c_{j-1}/x.c_{j+1}$ (if they are not leaves) to $x.c_j$. Then call "DELETE-NON-MINIMUM($x.c_j, k$)". If both $x.c_{j-1}.n$ and $x.c_{j+1} = t - 1$, merge $x.c_j$ with $x.c_{j-1}$, creating a new node with the same height as $x.c_j$, and delete k from the new node.

If k is at a leaf, DELETE-NON-MINIMUM is called h times, while at every level at most one merge would happen and takes constant time; otherwise, as searching for successor or predecessor would be needed, which takes $O(h)$ time.

(d) The input of this problem are a key k and two B-trees T_1 and T_2 with same minimum degree parameter t and height h_1 and h_2 . All keys in T_1 are strictly smaller than k while all keys in T_2 are strictly larger than k . The out put should be another B-tree with minimum degree parameter t and contains exactly all keys in T_1 and T_2 plus k

Let's define "COMBINE(r_1, r_2, k)", which takes two augmented B-tree root r_1 , and r_2 , and k , given that all keys in subtree below r_1 is strictly smaller than k and k is strictly smaller than all keys in subtree below r_2 , and both $r_1.n$ and $r_2.n < 2t - 1$. If $r_1.h = r_2.h$, do as (a); if $r_1.h > r_2.h$, we find the largest node in the subtree below r_1 with height $r_2.h$, and let it be r'_1 . Then combine r'_1 and r_2 , using the algorithm in (a). If the result, r_r has height $r_2.h$, replace r'_1 with r_r , and return r_1 . If r_r has height $r_2.h + 1$, add the only key in r_r to r_1 's parent node, and add the two children of r_r at the same time (thus overwrite r'_1). Now if r'_1 's parent has more than $2t - 1$ keys, split it and add a key to its parent, then split that node if it has $2t$ keys, ..., until the tree below r_1 is a valid B-tree. (Note that r_1 itself might be split, if so we let r_1 be its parent, namely, the new root node with only one key). Return r_1 . For $r_1.h < r_2.h$, the approach is the almost the same.

When $r_1.h > r_2.h$, finding r'_1 doesn't make any change to both trees. According to (a), the result of combining r'_1 and r_2 is a valid B-tree which contains all keys from r'_1 , r_2 and k . If $r_r.h = r_2.h$, then $r_r.n \leq t - 1$, replacing r'_1 with r_r makes r_1 a valid B-tree with all keys needed. If $r_r.h$ is one larger than $r_2.h$, then r_r should be added to a higher level. After the adding, r_1 is already a valid B-tree except for one node, which might have $2t$ keys. Splitting upward from that node, we can get a valid B-tree, which in this process the set of all keys under r_1 is not changed. The same for $r_1.h < r_2.h$.

Searching for r'_1 takes $O(|h_1 - h_2|)$ time (if $r_1.h \neq r_2.h$), combining r'_1 with r_2 takes constant time as in (a), and splitting $O(|h_1 - h_2|)$ times takes $O(|h_1 - h_2|)$ time. Considering when $h_1 = h_2$ constant time is needed, the overall time needed is $O(|h_1 - h_2| + 1)$.